

```
In [1]: import pandas as pd
import numpy as np

In [2]: df = pd.read_csv('/Users/sameerkhan/Documents/DATA CLEANING/Automobile price data _Raw_(before).csv')

In [3]: df
```

?	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...	130	mpfi	3.47	2.68	9.0	111	5000	21
?	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...	130	mpfi	3.47	2.68	9.0	111	5000	21
?	alfa-romero	gas	std	two	hatchback	rwd	front	94.5	...	152	mpfi	2.68	3.47	9.0	154	5000	19
164	audi	gas	std	four	sedan	fwd	front	99.8	...	109	mpfi	3.19	3.40	10.0	102	5500	24
164	audi	gas	std	four	sedan	4wd	front	99.4	...	136	mpfi	3.19	3.40	8.0	115	5500	18
...
95	volvo	gas	std	four	sedan	rwd	front	109.1	...	141	mpfi	3.78	3.15	9.5	114	5400	23
95	volvo	gas	turbo	four	sedan	rwd	front	109.1	...	141	mpfi	3.78	3.15	8.7	160	5300	19
95	volvo	gas	std	four	sedan	rwd	front	109.1	...	173	mpfi	3.58	2.87	8.8	134	5500	18
95	volvo	diesel	turbo	four	sedan	rwd	front	109.1	...	145	idi	3.01	3.40	23.0	106	4800	26

```
In [50]: df.shape

Out[50]: (205, 26)

In [4]: '''
1. Check for errors in datasets
2. Replace them with a suitable approach and explain why you opted for that method or technique.
3. Final report on the dataset before and after cleaning.
4. The final dataset needs to be saved to .csv format.
'''
df.describe()
```

Out[4]:

	symboling	wheel-base	length	width	height	curb-weight	engine-size	compression-ratio	city-mpg	highway-mpg
count	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000
mean	0.834146	98.756585	174.049268	65.907805	53.724878	2555.565854	126.907317	10.142537	25.219512	30.751220
std	1.245307	6.021776	12.337289	2.145204	2.443522	520.680204	41.642693	3.972040	6.542142	6.886443
min	-2.000000	86.600000	141.100000	60.300000	47.800000	1488.000000	61.000000	7.000000	13.000000	16.000000
25%	0.000000	94.500000	166.300000	64.100000	52.000000	2145.000000	97.000000	8.600000	19.000000	25.000000
50%	1.000000	97.000000	173.200000	65.500000	54.100000	2414.000000	120.000000	9.000000	24.000000	30.000000
75%	2.000000	102.400000	183.100000	66.900000	55.500000	2935.000000	141.000000	9.400000	30.000000	34.000000
max	3.000000	120.900000	208.100000	72.300000	59.800000	4066.000000	326.000000	23.000000	49.000000	54.000000

```
In [ ]: # cleaning 4wd errors and checking unique values to catogorical coumns

In [5]: import re

df.loc[df['drive-wheels']=='4wd','drive-wheels'] = 'fwd'
```

In [6]:

df

Out[6]:

name	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	...	engine-size	fuel-system	bore	stroke	compression-ratio	horsepower	peak-rpm	city-mpg
alfa-nero	gas	std	two	convertible	rwd	front	88.6	...	130	mpfi	3.47	2.68	9.0	111	5000	21
alfa-nero	gas	std	two	convertible	rwd	front	88.6	...	130	mpfi	3.47	2.68	9.0	111	5000	21
alfa-nero	gas	std	two	hatchback	rwd	front	94.5	...	152	mpfi	2.68	3.47	9.0	154	5000	19
audi	gas	std	four	sedan	fwd	front	99.8	...	109	mpfi	3.19	3.40	10.0	102	5500	24
audi	gas	std	four	sedan	fwd	front	99.4	...	136	mpfi	3.19	3.40	8.0	115	5500	18
...
volvo	gas	std	four	sedan	rwd	front	109.1	...	141	mpfi	3.78	3.15	9.5	114	5400	20
volvo	gas	turbo	four	sedan	rwd	front	109.1	...	141	mpfi	3.78	3.15	8.7	160	5300	19
volvo	gas	std	four	sedan	rwd	front	109.1	...	173	mpfi	3.58	2.87	8.8	134	5500	18
volvo	diesel	turbo	four	sedan	rwd	front	109.1	...	145	idi	3.01	3.40	23.0	106	4800	26
volvo	gas	turbo	four	sedan	rwd	front	109.1	...	141	mpfi	3.78	3.15	9.5	114	5400	19

In [10]:

df['normalized-losses'].unique()

Out[10]:

array(['?', '164', '158', '192', '188', '121', '98', '81', '118', '148', '110', '145', '137', '101', '78', '106', '85', '107', '104', '113', '150', '129', '115', '93', '142', '161', '153', '125', '128', '122', '103', '168', '108', '194', '231', '119', '154', '74', '186', '83', '102', '89', '87', '77', '91', '134', '65', '197', '90', '94', '256', '95'], dtype=object)

In [11]:

df['num-of-doors'].unique()

Out[11]:

array(['two', 'four', '?'], dtype=object)

In [12]:

df['bore'].unique()

Out[12]:

array(['3.47', '2.68', '3.19', '3.13', '3.50', '3.31', '3.62', '2.91', '3.03', '2.97', '3.34', '3.60', '2.92', '3.15', '3.43', '3.63', '3.54', '3.08', '?', '3.39', '3.76', '3.58', '3.46', '3.80', '3.78', '3.17', '3.35', '3.59', '2.99', '3.33', '3.70', '3.61', '3.94', '3.74', '2.54', '3.05', '3.27', '3.24', '3.01'], dtype=object)

In [13]:

df['stroke'].unique()

Out[13]:

array(['2.68', '3.47', '3.40', '2.80', '3.19', '3.39', '3.03', '3.11', '3.23', '3.46', '3.90', '3.41', '3.07', '3.58', '4.17', '2.76', '3.15', '?', '3.16', '3.64', '3.10', '3.35', '3.12', '3.86', '3.29', '3.27', '3.52', '2.19', '3.21', '2.90', '2.07', '2.36', '2.64', '3.08', '3.50', '3.54', '2.87'], dtype=object)

In [14]:

df['horsepower'].unique()

Out[14]:

array(['111', '154', '102', '115', '110', '140', '160', '101', '121', '182', '48', '70', '68', '88', '145', '58', '76', '60', '86', '100', '78', '90', '176', '262', '135', '84', '64', '120', '72', '123', '155', '184', '175', '116', '69', '55', '97', '152', '200', '95', '142', '143', '207', '288', '?', '73', '82', '94', '62', '56', '112', '92', '161', '156', '52', '85', '114', '162', '134', '106'], dtype=object)

In [15]:

df['peak-rpm'].unique()

Out[15]:

array(['5000', '5500', '5800', '4250', '5400', '5100', '4800', '6000', '4750', '4650', '4200', '4350', '4500', '5200', '4150', '5600', '5900', '5750', '?', '5250', '4900', '4400', '6600', '5300'], dtype=object)

```
In [22]: type(df['price'].value_counts())
```

```
Out[22]: pandas.core.series.Series
```

```
In [24]: df['stroke'] = pd.to_numeric(df['stroke'], errors = 'coerce')  
#converting it to numeric values from nan
```

```
In [27]: df[['stroke']].mode()  
#checking mode as its categorical column cant use mean to fill null values  
# as it will disturb the data  
df[['stroke']].mode().iloc[0]
```

```
Out[27]: stroke      3.4  
Name: 0, dtype: float64
```

```
In [28]: #replacing nan values with mode of column  
df['stroke'].replace(to_replace=np.nan,value=df['stroke'].mode()[0],inplace=True)
```

```
In [30]: df[['stroke']].isnull().sum()
```

```
Out[30]: stroke      0  
dtype: int64
```

```
In [31]: #same with other  
df['normalized-losses'] = pd.to_numeric(df['normalized-losses'], errors = 'coerce')
```

```
In [34]: df['normalized-losses'].mode().iloc[0]
```

```
Out[34]: 161.0
```

```
In [35]: df['normalized-losses'].replace(to_replace=np.nan,value=df['normalized-losses'].mode()[0],inplace=True)
```

```
In [36]: df['bore'] = pd.to_numeric(df['bore'], errors = 'coerce')
```

```
In [38]: df['bore'].mode().iloc[0]
```

```
Out[38]: 3.62
```

```
In [39]: df['bore'].replace(to_replace=np.nan,value=df['bore'].mode()[0],inplace=True)
```

```
In [43]: df['horsepower'].mode()
```

```
Out[43]: 0      68.0  
dtype: float64
```

```
In [41]: df['horsepower'] = pd.to_numeric(df['horsepower'], errors = 'coerce')
```

```
In [42]: df['horsepower'].replace(to_replace=np.nan,value=df['horsepower'].mode()[0],inplace=True)
```

```
In [44]: df['peak-rpm'] = pd.to_numeric(df['peak-rpm'], errors = 'coerce')
```

```
In [46]: df['peak-rpm'].mode()
```

```
Out[46]: 0      5500.0  
dtype: float64
```

```
In [47]: df['peak-rpm'].replace(to_replace=np.nan,value=df['peak-rpm'].mode()[0],inplace=True)
```

```
In [48]: df['num-of-doors'].mode()
```

```
Out[48]: 0      four  
dtype: object
```

```
In [49]: df['num-of-doors'].replace(to_replace='?',value=df['num-of-doors'].mode()[0],inplace=True)
```

In [59]: `df.info()`
`df.shape`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  -
0   symboling              205 non-null    int64
1   normalized-losses      205 non-null    float64
2   make                   205 non-null    object
3   fuel-type              205 non-null    object
4   aspiration              205 non-null    object
5   num-of-doors            205 non-null    object
6   body-style              205 non-null    object
7   drive-wheels            205 non-null    object
8   engine-location         205 non-null    object
9   wheel-base              205 non-null    float64
10  length                  205 non-null    float64
11  width                   205 non-null    float64
12  height                  205 non-null    float64
13  curb-weight             205 non-null    int64
14  engine-type             205 non-null    object
15  num-of-cylinders        205 non-null    int64
16  engine-size             205 non-null    float64
17  fuel-system             205 non-null    object
18  bore                    205 non-null    float64
19  stroke                  205 non-null    float64
20  compression-ratio       205 non-null    float64
21  horsepower              205 non-null    int64
22  peak-rpm                205 non-null    int64
23  city-mpg                 205 non-null    int64
24  highway-mpg              205 non-null    int64
25  price                   205 non-null    int64
dtypes: int64(10), float64(10), object(6)
```

In [63]: `df.isna().sum()`

```
drive-wheels      0
engine-location    0
wheel-base        0
length            0
width             0
height            0
curb-weight        0
engine-type        0
num-of-cylinders   0
engine-size        0
fuel-system        0
bore              0
stroke            0
compression-ratio  0
horsepower         0
peak-rpm          0
city-mpg           0
highway-mpg        0
price             0
dtypes: int64(10), float64(10), object(6)
```

In [66]: `df.to_csv('/Users/sameerkhan/Documents/DATA CLEANING/Auto_mobile_dataset_cleaned1.csv')`

In [68]: `df['price'].unique()`

```
'12945', '10345', '6785', '11048', '32250', '35550', '36000',
'5195', '6095', '6795', '6695', '7395', '10945', '11845', '13645',
'15645', '8495', '10595', '10245', '10795', '11245', '18280',
'18344', '25552', '28248', '28176', '31600', '34184', '35056',
'40960', '45400', '16503', '5389', '6189', '6669', '7689', '9959',
'8499', '12629', '14869', '14489', '6989', '8189', '9279', '5499',
'7099', '6649', '6849', '7349', '7299', '7799', '7499', '7999',
'8249', '8949', '9549', '13499', '14399', '17199', '19699',
'18399', '11900', '13200', '12440', '13860', '15580', '16900',
'16695', '17075', '16630', '17950', '18150', '12764', '22018',
'32528', '34028', '37028', '9295', '9895', '11850', '12170',
'15040', '15510', '18620', '5118', '7053', '7603', '7126', '7775',
'9960', '9233', '11259', '7463', '10198', '8013', '11694', '5348',
'6338', '6488', '6918', '7898', '8778', '6938', '7198', '7788',
'7738', '8358', '9258', '8058', '8238', '9298', '9538', '8449',
'9639', '9989', '11199', '11549', '17669', '8948', '10698', '9988',
'10898', '11248', '16558', '15998', '15690', '15750', '7975',
'7995', '8195', '9495', '9995', '11595', '9980', '13295', '13845',
'12290', '12940', '13415', '15985', '16515', '18420', '18950',
'16845', '10045', '21485', '22470', '22625', dtype=object)
```

In [72]: `df['price'] = pd.to_numeric(df['price'], errors = 'coerce')`

In [76]: `df['price'].mean()`

Out[76]: 13207.129353233831

```
In [77]: df.isna().sum()
```

```
Out[77]: symboling      0
normalized-losses    0
make                 0
fuel-type            0
aspiration           0
num-of-doors         0
body-style           0
drive-wheels         0
engine-location      0
wheel-base          0
length              0
width               0
height              0
curb-weight          0
engine-type          0
num-of-cylinders     0
engine-size          0
fuel-system          0
bore                 0
stroke              0
```

```
In [78]: df['price'].isna().sum()
```

Out[78]: 4

```
In [79]: df['price'].replace(to_replace=np.nan,value=df['price'].mean(),inplace=True)
```

```
In [80]: df['price'].isna().sum()
```

Out[80]: 0

```
In [81]: df.to_csv('/Users/sameerkhan/Documents/DATA CLEANING/Auto_mobile_dataset_cleaned_final.csv')
```

```
In [82]: df.shape
```

Out[82]: (205, 26)

```
In [ ]:
```